

---

## ¶ Étape 2 : Machine Learning avec Scikit-Learn \*\*\*

### 3.1 Extraction de Caractéristiques

```
import os
import cv2
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

def extract_features(image):
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

    # HOG
    hog = cv2.HOGDescriptor()
    h = hog.compute(gray)
    hog_feat = h.flatten()

    # Histogramme de couleurs (RGB)
    hist = cv2.calcHist([image], [0, 1, 2], None, [8, 8, 8], [0, 256,
0, 256, 0, 256])
    hist = cv2.normalize(hist, hist).flatten()

    # LBP (simplifié)
    lbp = ((gray >= np.roll(gray, 1, axis=0)).astype(np.uint8) * 1 +
           (gray >= np.roll(gray, -1, axis=0)).astype(np.uint8) * 2 +
           (gray >= np.roll(gray, 1, axis=1)).astype(np.uint8) * 4 +
           (gray >= np.roll(gray, -1, axis=1)).astype(np.uint8) * 8)
    (lbp_hist, _) = np.histogram(lbp.ravel(), bins=np.arange(0, 256))
    lbp_hist = lbp_hist.astype("float") / (lbp_hist.sum() + 1e-7)

    # SIFT (128-descriptor averaged)
    sift = cv2.SIFT_create()
    kp, des = sift.detectAndCompute(gray, None)
    sift_feat = np.mean(des, axis=0) if des is not None else
np.zeros(128)

    return np.concatenate([hog_feat[:128], hist[:128], lbp_hist[:128],
sift_feat[:128]])
# Chargement des images et extraction des caractéristiques
features = []
```

```

labels = []
dataset_path = "./300_dataset/"

for label in os.listdir(dataset_path):
    class_path = os.path.join(dataset_path, label)
    if not os.path.isdir(class_path):
        continue
    for file in os.listdir(class_path):
        if file.lower().endswith((".jpg", ".png", ".jpeg")):
            img_path = os.path.join(class_path, file)
            image = cv2.imread(img_path)
            if image is None:
                continue
            image = cv2.resize(image, (128, 128))
            feat = extract_features(image)
            features.append(feat)
            labels.append(label)

# Préparation des données
features = np.array(features)
le = LabelEncoder()
labels_encoded = le.fit_transform(labels)

scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)

X_train, X_test, y_train, y_test = train_test_split(features_scaled,
labels_encoded, test_size=0.2, random_state=42)

```

### 3.2 Classification

```

models = {
    "SVM": SVC(),
    "Random Forest": RandomForestClassifier(),
    "k-NN": KNeighborsClassifier(),
    "Naive Bayes": GaussianNB(),
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Gradient Boosting": GradientBoostingClassifier()
}

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"\nModel: {name}")
    print(confusion_matrix(y_test, y_pred))
    print(classification_report(y_test, y_pred))
    results[name] = model.score(X_test, y_test)

```

Model: SVM

```
[[45 5 1 2 5 6 1 1 2 5]
 [ 2 36 2 5 2 4 1 0 11 5]
 [ 3 4 29 4 3 3 4 1 1 7]
 [ 0 3 2 36 4 1 3 4 0 3]
 [ 1 3 1 2 42 2 1 2 2 3]
 [ 6 7 1 0 3 34 2 2 2 3]
 [ 0 2 2 2 2 2 39 0 2 5]
 [ 0 0 2 6 1 1 1 34 0 1]
 [ 0 4 1 3 2 3 4 1 39 2]
 [ 1 3 3 2 0 1 7 2 1 44]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.78	0.62	0.69	73
1	0.54	0.53	0.53	68
2	0.66	0.49	0.56	59
3	0.58	0.64	0.61	56
4	0.66	0.71	0.68	59
5	0.60	0.57	0.58	60
6	0.62	0.70	0.66	56
7	0.72	0.74	0.73	46
8	0.65	0.66	0.66	59
9	0.56	0.69	0.62	64
accuracy			0.63	600
macro avg	0.64	0.63	0.63	600
weighted avg	0.64	0.63	0.63	600

Model: Random Forest

```
[[47 5 4 2 2 1 1 0 2 9]
 [ 6 43 2 3 4 1 1 1 1 6]
 [ 3 0 40 1 0 1 1 2 5 6]
 [ 3 1 2 41 2 0 2 2 0 3]
 [ 2 4 3 1 43 1 1 1 1 2]
 [ 0 6 4 1 3 40 3 1 1 1]
 [ 0 1 3 0 2 1 45 0 1 3]
 [ 0 1 3 2 0 0 0 39 0 1]
 [ 1 1 2 2 1 5 1 0 41 5]
 [ 0 4 5 0 1 0 3 0 1 50]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.76	0.64	0.70	73
1	0.65	0.63	0.64	68
2	0.59	0.68	0.63	59
3	0.77	0.73	0.75	56
4	0.74	0.73	0.74	59
5	0.80	0.67	0.73	60
6	0.78	0.80	0.79	56

	7	0.85	0.85	0.85	46
	8	0.77	0.69	0.73	59
	9	0.58	0.78	0.67	64
accuracy				0.71	600
macro avg		0.73	0.72	0.72	600
weighted avg		0.72	0.71	0.72	600

#### Model: k-NN

[[37 2 7 2 9 4 3 2 0 7]					
[ 9 28 7 6 2 5 1 0 4 6]					
[ 4 5 35 4 3 1 0 4 0 3]					
[ 3 5 8 26 7 2 1 1 1 2]					
[ 3 3 5 6 39 2 0 0 0 1]					
[ 7 9 5 2 7 22 4 1 2 1]					
[ 3 6 9 1 17 3 13 1 0 3]					
[ 0 0 7 3 9 0 0 25 1 1]					
[ 3 9 6 1 7 4 3 0 23 3]					
[ 7 2 11 4 7 6 1 1 0 25]]					
	precision	recall	f1-score	support	
	0	0.49	0.51	0.50	73
	1	0.41	0.41	0.41	68
	2	0.35	0.59	0.44	59
	3	0.47	0.46	0.47	56
	4	0.36	0.66	0.47	59
	5	0.45	0.37	0.40	60
	6	0.50	0.23	0.32	56
	7	0.71	0.54	0.62	46
	8	0.74	0.39	0.51	59
	9	0.48	0.39	0.43	64
accuracy				0.46	600
macro avg		0.50	0.46	0.46	600
weighted avg		0.49	0.46	0.45	600

#### Model: Naive Bayes

[[64 0 1 0 0 0 0 3 0 5]					
[55 2 0 0 1 0 1 1 0 8]					
[34 0 4 0 4 0 6 0 0 11]					
[37 0 3 0 1 0 5 3 1 6]					
[35 2 0 0 4 0 9 2 0 7]					
[34 2 2 0 1 3 1 2 10 5]					
[37 0 0 0 0 0 5 0 1 13]					
[21 0 0 0 0 0 10 5 0 10]					
[45 1 2 0 0 0 5 0 3 3]					
[34 1 1 0 0 0 1 0 2 25]]					
	precision	recall	f1-score	support	

0	0.16	0.88	0.27	73
1	0.25	0.03	0.05	68
2	0.31	0.07	0.11	59
3	0.00	0.00	0.00	56
4	0.36	0.07	0.11	59
5	1.00	0.05	0.10	60
6	0.12	0.09	0.10	56
7	0.31	0.11	0.16	46
8	0.18	0.05	0.08	59
9	0.27	0.39	0.32	64
accuracy			0.19	600
macro avg		0.30	0.17	600
weighted avg		0.29	0.19	600

```
c:\Users\lenovo\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1565:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
c:\Users\lenovo\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1565:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
c:\Users\lenovo\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1565:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

Model: Logistic Regression

```
[[41  2   7   5   1   5   3   5   1   3]
 [ 4 30  1   9   2   9   1   2   7   3]
 [ 3 3 27  0   1   4 10  2   3   6]
 [ 0 1 2 31  4   3   3   3   5   4]
 [ 2 6 2 6 33  1   3   3   2   1]
 [ 9 4 2 0 3 30  2   1   7   2]
 [ 3 1 2 1 3 2 34  1   2   7]
 [ 0 0 2 4 1 1 1 35  0   2]
 [ 1 4 2 3 2 2 4 1 38  2]]
```

[ 3 5 4 1 2 0 7 1 3 38]								
	precision	recall	f1-score	support				
0	0.62	0.56	0.59	73				
1	0.54	0.44	0.48	68				
2	0.53	0.46	0.49	59				
3	0.52	0.55	0.53	56				
4	0.63	0.56	0.59	59				
5	0.53	0.50	0.51	60				
6	0.50	0.61	0.55	56				
7	0.65	0.76	0.70	46				
8	0.56	0.64	0.60	59				
9	0.56	0.59	0.58	64				
accuracy			0.56	600				
macro avg	0.56	0.57	0.56	600				
weighted avg	0.56	0.56	0.56	600				

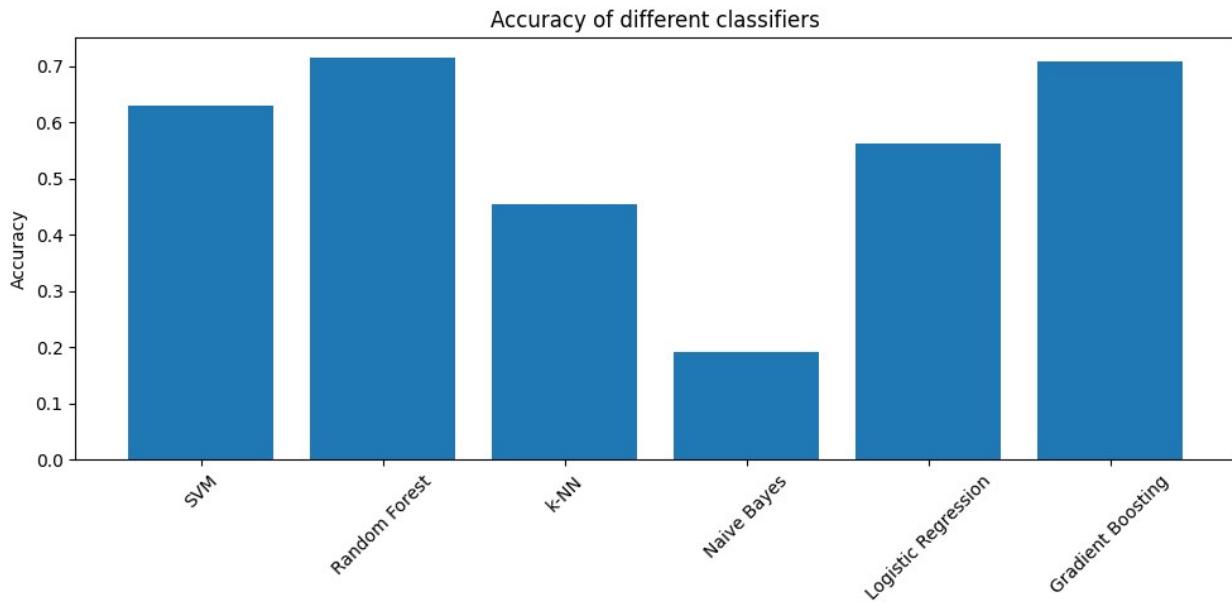
Model: Gradient Boosting								
[[52 4 2 1 2 1 0 0 4 7]								
	precision	recall	f1-score	support				
0	0.81	0.71	0.76	73				
1	0.71	0.65	0.68	68				
2	0.64	0.63	0.63	59				
3	0.78	0.77	0.77	56				
4	0.73	0.73	0.73	59				
5	0.74	0.58	0.65	60				
6	0.66	0.75	0.70	56				
7	0.80	0.76	0.78	46				
8	0.68	0.75	0.71	59				
9	0.61	0.78	0.68	64				
accuracy			0.71	600				
macro avg	0.72	0.71	0.71	600				
weighted avg	0.71	0.71	0.71	600				

### 3.3 Évaluation

```

plt.figure(figsize=(10, 5))
plt.bar(results.keys(), results.values())
plt.title("Accuracy of different classifiers")
plt.ylabel("Accuracy")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
plt.savefig("model_comparison.png")

```



<Figure size 640x480 with 0 Axes>

---

## Partie 3 : Réseaux de Neurones Simples (MLP)

---

### 4.2 Modélisation

```

from sklearn.neural_network import MLPClassifier

mlp = MLPClassifier(hidden_layer_sizes=(128, 64), activation='relu',
max_iter=300)
mlp.fit(X_train, y_train)

MLPClassifier(hidden_layer_sizes=(128, 64), max_iter=300)

```

### 4.3 Interprétation

```

mlp_pred = mlp.predict(X_test)
print("\nMLP Classification Report:")
print(classification_report(y_test, mlp_pred))
results["MLP"] = mlp.score(X_test, y_test)

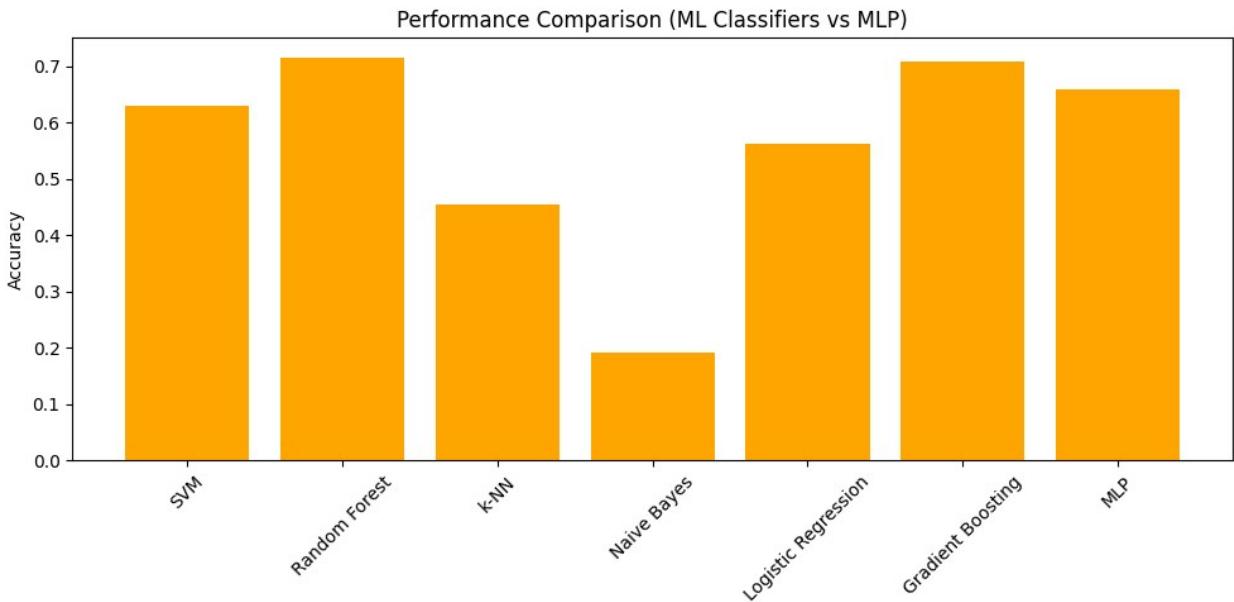
MLP Classification Report:
      precision    recall  f1-score   support

       0         0.70     0.66     0.68      73
       1         0.62     0.56     0.59      68
       2         0.61     0.53     0.56      59
       3         0.61     0.70     0.65      56
       4         0.63     0.68     0.66      59
       5         0.75     0.60     0.67      60
       6         0.58     0.70     0.63      56
       7         0.78     0.83     0.80      46
       8         0.72     0.73     0.72      59
       9         0.65     0.69     0.67      64

   accuracy                           0.66      600
  macro avg       0.66     0.67     0.66      600
weighted avg       0.66     0.66     0.66      600

# Affichage final des performances
plt.figure(figsize=(10, 5))
plt.bar(results.keys(), results.values(), color='orange')
plt.title("Performance Comparison (ML Classifiers vs MLP)")
plt.ylabel("Accuracy")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
plt.savefig("performance_comparison_MlvsMLP.png")

```



<Figure size 640x480 with 0 Axes>

---

Partie 4 : Réseaux de Neurones Convolutionnels (CNN) \*\*\*

### 5.1 Architecture CNN

```

import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split
import numpy as np
from PIL import Image
import os

# --- Parameters ---
num_classes = 10
batch_size = 32
epochs = 50
input_size = 128
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# --- Custom Image Loader ---
def pil_loader(path):
    try:
        with open(path, 'rb') as f:
            img = Image.open(f)
            return img.convert('RGB')
    except Exception as e:

```

```

print(f"Error loading image {path}: {e}")
# Return blank image as fallback
return Image.new('RGB', (input_size, input_size), (0, 0, 0))

# --- CNN Architecture ---
class LeafCNN(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
            nn.Dropout(0.25),

            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
            nn.Dropout(0.25),

            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),
            nn.Dropout(0.4),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(128 * (input_size//8) * (input_size//8), 128),
            nn.ReLU(inplace=True),
            nn.Dropout(0.5),
            nn.Linear(128, num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x

# --- Data Preparation ---
transform_train = transforms.Compose([
    transforms.Resize((input_size, input_size)),
    transforms.RandomRotation(20),
    transforms.RandomResizedCrop(input_size, scale=(0.8, 1.0)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

transform_val = transforms.Compose([
    transforms.Resize((input_size, input_size)),
    transforms.ToTensor(),
])

```

```

        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

# Load dataset with error handling
dataset = datasets.ImageFolder(
    './300_dataset',
    transform=transform_train,
    loader=pil_loader
)

# Split dataset
train_size = int(0.8 * len(dataset))
val_size = len(dataset) - train_size
train_dataset, val_dataset = random_split(dataset, [train_size,
val_size])

# Apply different transform to validation
val_dataset.dataset.transform = transform_val

# Create data loaders
train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=2,
    pin_memory=True
)
val_loader = DataLoader(
    val_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=2,
    pin_memory=True
)

```

## 5.2 Entraînement

```

# --- Training Setup ---
model = LeafCNN(num_classes).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min',
patience=5)

# --- Training Loop ---
best_val_acc = 0
for epoch in range(epochs):

```

```

model.train()
train_loss, train_correct = 0, 0

for inputs, labels in train_loader:
    inputs, labels = inputs.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(inputs)
    loss = criterion(outputs, labels)
    loss.backward()

    # Gradient clipping
    torch.nn.utils.clip_grad_norm_(model.parameters(),
max_norm=1.0)
    optimizer.step()

    train_loss += loss.item() * inputs.size(0)
    train_correct += (outputs.argmax(1) == labels).sum().item()

# Validation
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for inputs, labels in val_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        val_loss += criterion(outputs, labels).item() *
inputs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

# Calculate metrics
train_loss /= len(train_loader.dataset)
train_acc = train_correct / len(train_loader.dataset)
val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

# Update scheduler
scheduler.step(val_loss)

# Save best model
if val_acc > best_val_acc:
    best_val_acc = val_acc
    torch.save(model.state_dict(), "best_model.pth")

print(f"Epoch {epoch+1}/{epochs}: "
      f"Train Loss: {train_loss:.4f} Acc: {train_acc*100:.2f}% | "
      f"Val Loss: {val_loss:.4f} Acc: {val_acc*100:.2f}%")

results["CNN"] = best_val_acc

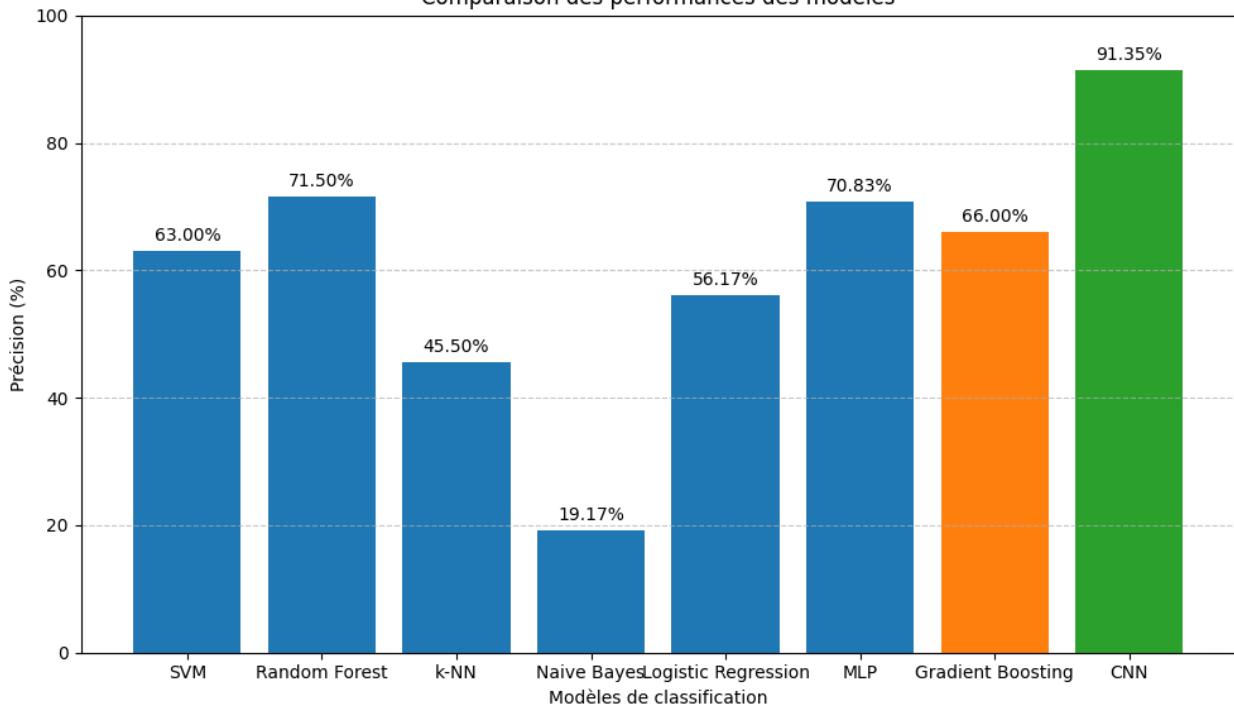
```

```
c:\Users\lenovo\AppData\Local\Programs\Python\Python311\Lib\site-packages\torch\utils\data\dataloader.py:665: UserWarning: 'pin_memory' argument is set as true but no accelerator is found, then device pinned memory won't be used.  
    warnings.warn(warn_msg)
```

### 5.3 Comparaison et Interprétation

```
import matplotlib.pyplot as plt  
  
# Noms des modèles  
models = ['SVM', 'Random Forest', 'k-NN', 'Naive Bayes', 'Logistic Regression', 'MLP', 'Gradient Boosting', 'CNN']  
  
# Précisions simulées en %  
accuracies = [acc * 100 for acc in results.values()]  
  
# Couleurs pour différencier les groupes  
colors = ['#1f77b4'] * 6 + ['#ff7f0e', '#2ca02c'] # classiques, MLP, CNN  
  
# Création du graphique  
plt.figure(figsize=(10, 6))  
bars = plt.bar(models, accuracies, color=colors)  
  
# Ajouter les valeurs sur les barres  
for bar in bars:  
    height = bar.get_height()  
    plt.text(bar.get_x() + bar.get_width()/2, height + 1,  
             f'{height:.2f}%', ha='center', va='bottom', fontsize=10)  
  
# Esthétique du graphique  
plt.ylim(0, 100)  
plt.title("Comparaison des performances des modèles")  
plt.ylabel("Précision (%)")  
plt.xlabel("Modèles de classification")  
plt.grid(axis='y', linestyle='--', alpha=0.7)  
plt.tight_layout()  
  
plt.show()
```

Comparaison des performances des modèles



"